

DIFFERENTIAL COUNT ANALYSIS WITH A TOPIC MODEL

PETER CARBONETTO*

1. Motivation, and overview of methods. The aim of this document is to derive, from first principles, a method for analysis of differential gene expression using a topic model (also known as “grade of membership model” [3]). This method may have other uses — say, to identify “key words” in a topic modeling analysis of text documents — but since our main motivation is analysis of gene expression data, we describe the methods with that application in mind.

For motivation, we begin with the “log-fold change” statistic commonly used in microarray and RNA sequencing experiments to quantify expression differences between two conditions (e.g., [2, 4]). The log-fold change for gene j and condition k is a ratio of two conditional expectations,

$$(1) \quad \text{lfc}(j, k) \equiv \log_2 \frac{E[x_j \mid \text{condition} = k]}{E[x_j \mid \text{condition} \neq k]},$$

where x_j is the measured expression level (e.g., UMI count) of gene j .¹

The statistic (1) is a measure of *absolute* change in expression level between two conditions. It is often preferable to measure *relative* change, say, change in gene expression relative to the total expression in each sample. As we will see, a topic modeling perspective provides a natural way to analyze both absolute or relative change in gene expression.

Topic modeling brings two new twists to analysis of differential gene expression:

1. In conventional differential expression analysis, each sample is assigned to a single condition; in topic modeling, the topic assignments are proportional.
2. In conventional differential expression analysis, all genes expression measurements in a sample are assumed to be from the same condition; in topic modeling, the topic assignments are made at the level of genes, not samples.

We will need to bear in mind these two points in developing the models and methods.

1.1. The binomial model. We begin with a simple binomial model of gene expression given the proportional assignments to a topic,

$$(2) \quad x_i \sim \text{Binom}(n_i, \pi_i),$$

where x_i is the expression level of the gene in sample i , n_i is the total expression in sample i , and $\text{Binom}(n, \theta)$ denotes the binomial distribution with n trials and success probability θ . In this model, the binomial probabilities are defined as

$$(3) \quad \pi_i = (1 - q_i)f_0 + q_i f_1,$$

where $q_i \in [0, 1]$ is the known proportion of sample i attributed to the topic, and $f_0, f_1 \in [0, 1]$ are two unknowns to be estimated.

*Dept. of Human Genetics and the Research Computing Center, University of Chicago, Chicago, IL

¹Defining x_{jk} as the total gene expression for gene j among all samples (expression profiles) in condition k , x_j as the total gene expression for gene j in all samples, n_k as the number of samples in condition k , and n as the total sample size, the log-fold change can be computed as $\text{lfc}(j, k) = \log_2 \left\{ \frac{x_{jk}}{x_j - x_{jk}} \times \frac{n - n_k}{n_k} \right\}$.

In econometrics, this model is sometimes called a *linear probability model* [5]. This viewpoint is more apparent by rewriting the binomial probabilities as

$$(4) \quad \pi_i = \beta_0 + q_i \beta,$$

where $\beta_0 = f_0$ and $\beta = f_1 - f_0 \in [-1, +1]$. This is a simple regression model for the binomial probability, in which the binomial probability increases linearly with topic proportion q_i , with the slope given by β .

Next, we show that this simple binomial model implements differential gene expression: specifically, $\log_2(f_1/f_0)$ quantifies relative log-fold change.

1.2. The Poisson model.

2. The multinomial topic model and Poisson non-negative matrix factorization. Here we briefly describe the multinomial topic model, and its connection to Poisson non-negative matrix factorization (Poisson NMF).

We begin with the “bag of words” description, which was used to describe the LDA model [1]. In this view, each document (or gene expression profile) i is represented as a vector of terms/genes, $w_i = (w_{i1}, \dots, w_{is_i})$, where s_i is the size of document i . (The order of the words or genes appearing in this vector doesn’t matter, hence the “bag of words.”) Each $w_{it} \in \{1, \dots, m\}$ is term/gene j with probability $p(w_{it} = j \mid z_{it} = k) = f_{jk}$, in which we have introduced z_{it} , a variable indicating which topic $k \in \{1, \dots, K\}$ the word/gene is drawn from. The topic indicator variables for document i are in turn generated according to $p(z_{it} = k) = l_{ik}$.

This process also defines a *multinomial* model for an $n \times m$ matrix of counts x_{ij} :

$$(5) \quad x_{i1}, \dots, x_{im} \sim \text{Multinom}(x_{i1}, \dots, x_{im}; s_i, \pi_i),$$

where $x_{ij} = \sum_{t=1}^{s_i} \delta_j(w_{it})$ is the number of times term/gene j appears in document/cell i , and the probabilities π_{ij} are weighted sums of the “factors” f_{jk} ,

$$(6) \quad \pi_{ij} = \sum_{k=1}^K l_{ik} f_{jk}.$$

The log-likelihood for the multinomial topic model, ignoring terms that do not depend on the model parameters, has a simple expression:

$$(7) \quad \log p(x) = \sum_{i=1}^n \sum_{j=1}^m x_{ij} \log(\sum_{k=1}^K l_{ik} f_{jk}).$$

As we have shown elsewhere, the multinomial topic model is closely related to a Poisson non-negative matrix factorization of the count data,

$$(8) \quad x_{ij} \sim \text{Poisson}(\lambda_{ij}),$$

where $\lambda_{ij} = \sum_{k=1}^K \hat{l}_{ik} \hat{f}_{jk}$. Given a Poisson NMF fit, an equivalent multinomial topic model can be easily recovered, as we have shown elsewhere.

3. Gene expression differences in topics. Returning to the question of assessing differential gene expression, there are two new twists when done in the context of topic modeling:

1. The cluster (topic) assignments are probabilistic.

2. The cluster assignments are made at the level of genes, not cells.

I propose a log-fold change statistic to address these two points. It compares the probability of gene j occurring ($w = j$) given topic k ($z = k$) versus the probability given assignment a topic other than k ($z \neq k$):

$$(9) \quad \text{lfc}^{\text{topics}}(j, k) \equiv \log_2 \frac{p(w = j \mid z = k)}{p(w = j \mid z \neq k)}.$$

For a given gene j and topic k , $\text{lfc}(j, k)$ can be calculated as

$$(10) \quad \begin{aligned} \text{lfc}^{\text{topics}}(j, k) &= \log_2 \left\{ \frac{p(w = j, z = k)}{p(w = j, z \neq k)} \times \frac{p(z = k)}{p(z \neq k)} \right\} \\ &= \log_2 \left\{ \frac{\sum_{i=1}^n \sum_{t=1}^{s_i} \delta_j(w_{it}) \phi_{ijkt}}{\sum_{i=1}^n \sum_{t=1}^{s_i} \delta_j(w_{it}) (1 - \phi_{ijkt})} \right. \\ &\quad \times \left. \frac{\sum_{i=1}^n \sum_{j'=1}^m \sum_{t=1}^{s_i} \delta_{j'}(w_{it}) (1 - \phi_{ij'kt})}{\sum_{i=1}^n \sum_{j'=1}^m \sum_{t=1}^{s_i} \delta_{j'}(w_{it}) \phi_{ij'kt}} \right\}, \end{aligned}$$

where ϕ_{ijkt} denotes the posterior probability of $z_{it} = k$ given $w_{it} = j$,

$$(11) \quad \begin{aligned} \phi_{ijkt} &\equiv p(z_{it} = k \mid w_{it} = j) \\ &= \frac{p(w_{it} = j \mid z_{it} = k) p(z_{it} = k)}{\sum_{k'=1}^K p(w_{it} = j \mid z_{it} = k') p(z_{it} = k')} \\ &= \frac{l_{ik} f_{jk}}{\sum_{k'=1}^K l_{ik'} f_{jk'}}. \end{aligned}$$

Since the topic assignments z_{it} do not depend on t —that is, we can drop the “ t ” subscript from the ϕ_{ijkt} ’s—the expression for the lfc simplifies:

$$(12) \quad \text{lfc}^{\text{topics}}(j, k) = \log_2 \left\{ \frac{\sum_{i=1}^n x_{ij} \phi_{ijk}}{\sum_{i=1}^n x_{ij} (1 - \phi_{ijk})} \times \frac{\sum_{i=1}^n \sum_{j'=1}^m x_{ij'} (1 - \phi_{ij'k})}{\sum_{i=1}^n \sum_{j'=1}^m x_{ij'} \phi_{ij'k}} \right\}.$$

At the maximum-likelihood solution (MLE) of the l_{ik} ’s and f_{jk} ’s, the lfc statistic simplifies further:

$$(13) \quad \text{lfc}^{\text{topics}}(j, k) = \log_2 \left\{ \frac{\sum_{i=1}^n x_{ij} \phi_{ijk}}{\sum_{i=1}^n x_{ij} (1 - \phi_{ijk})} \times \frac{\sum_{i=1}^n s_i (1 - l_{ik})}{\sum_{i=1}^n s_i l_{ik}} \right\}.$$

This is because, at the MLE, the loadings l_{ik} , $k = 1, \dots, K$, for a given document/cell i should be equal to the average of the weighted counts $\frac{1}{s_i} \sum_{j=1}^m x_{ij} \phi_{ijk}$.

Finally, it is convenient that the lfc (10, 13) will be the same if we replace the multinomial topic model parameters l_{ik} and f_{jk} with the corresponding parameters of the Poisson NMF, \hat{l}_{ik} and \hat{f}_{jk} (proof not given). From the derivation of the EM algorithm for Poisson NMF, this identity holds at the MLE:

$$\hat{f}_{jk} = \frac{\sum_{i=1}^n \phi_{ijk}}{\sum_{i=1}^n \hat{l}_{ik}}.$$

Plugging this relationship into (13), we obtain the following simple expression for the log-fold change:

$$(14) \quad \text{lfc}^{\text{topics}}(j, k) = \log_2 \left\{ \frac{\hat{f}_{jk} \sum_{i=1}^n \hat{l}_{ik}}{\sum_{k' \neq k} \hat{f}_{jk'} \sum_{i=1}^n \hat{l}_{ik'}} \times \frac{\sum_{i=1}^n s_i (1 - \hat{l}_{ik})}{\sum_{i=1}^n s_i \hat{l}_{ik}} \right\}.$$

What is nice about this about this expression is that it can be computed without seeing the data. It is also plain to see from this expression that to arrive at a log-fold change, one must weight the factors f_{jk} by the sample-wide topic probabilities $\sum_i l_{ik}$ across This same expression also works with the for the parameters of multinomial topic model l_{ik}, f_{jk} , again, so long as they are MLEs (proof not shown).

REFERENCES

- [1] D. M. BLEI, A. Y. NG, AND M. I. JORDAN, *Latent Dirichlet allocation*, Journal of Machine Learning Research, 3 (2003), pp. 993–1022.
- [2] X. CUI AND G. A. CHURCHILL, *Statistical tests for differential expression in cDNA microarray experiments*, Genome Biology, 4 (2003).
- [3] K. K. DEY, C. J. HSIAO, AND M. STEPHENS, *Visualizing the structure of RNA-seq expression data using grade of membership models*, PLoS Genetics, 13 (2017), p. e1006599.
- [4] J. QUACKENBUSH, *Microarray data normalization and transformation*, Nature Genetics, 32 (2002), pp. 496–501.
- [5] J. H. STOCK AND M. W. WATSON, *Introduction to Econometrics*, Pearson, Boston, MA, 3rd ed., 2015.