

SKETCH OF THE “ALTERNATING SQP” METHOD FOR FITTING POISSON TOPIC MODELS

PETER CARBONETTO*

1. Some derivations. Given an $n \times p$ matrix of counts X with entries $x_{ij} \geq 0$, the aim is to fit a Poisson model of the counts,

$$(1.1) \quad \begin{aligned} p(x) &= \prod_{i=1}^n \prod_{j=1}^p p(x_{ij}) \\ &= \prod_{i=1}^n \prod_{j=1}^p \text{Poisson}(x_{ij}; \lambda_{ij}), \end{aligned}$$

in which the Poisson rates are given by $\lambda_{ij} = \sum_{k=1}^K l_{ik} f_{jk}$. The model is determined by a $p \times K$ matrix F with entries $f_{ik} \geq 0$ (the “factors”) and an $n \times K$ matrix L with entries $l_{ik} \geq 0$ (the “loadings”). Fitting F and L is equivalent to non-negative matrix factorization with the “beta-divergence” cost function [3]. It can also be used to recover a maximum-likelihood estimate for the latent Dirichlet allocation (LDA) model [1]. So fitting this model is useful for a wide range of applications.

The log-likelihood for the Poisson model is

$$(1.2) \quad \log p(x | F, L) \propto \sum_{i=1}^n \sum_{j=1}^p x_{ij} \log(\sum_{k=1}^K l_{ik} f_{jk}) - \sum_{i=1}^n \sum_{j=1}^p \sum_{k=1}^K l_{ik} f_{jk},$$

where the constant of proportionality is obtained from factorial terms in the Poisson densities. Our specific aim is to find a F and L that maximizes the log-likelihood (1.2); that is, we would like to solve

$$(1.3) \quad \begin{aligned} &\text{minimize} && -\log p(x | F, L) \\ &\text{subject to} && F \geq 0, L \geq 0. \end{aligned}$$

In the remainder, we derive an efficient approach to doing this.

Our strategy for solving (1.3) is to alternate between solving for F with L fixed, and solving for L with F fixed. When solving for F with L fixed (and vice versa), the problem naturally decomposes into a collection of much smaller subproblems that are much more tractable to solve. All the subproblems are of the following form:

$$(1.4) \quad \begin{aligned} &\text{minimize} && \phi(y; B, w) \\ &\text{subject to} && y_k \geq 0 \text{ for all } k = 1, \dots, K, \end{aligned}$$

in which the objective function is defined as

$$(1.5) \quad \phi(y; B, w) = - \sum_{i=1}^n w_i \log \left(\sum_{k=1}^K b_{ik} y_k \right) + \sum_{i=1}^n \sum_{k=1}^K b_{ik} y_k.$$

*Dept. of Human Genetics and the Research Computing Center, University of Chicago, Chicago, IL

To see the connection between subproblem (1.4) and the original optimization problem (1.3), observe that the negative log-likelihood can be recovered as

$$(1.6) \quad -\log p(x | F, L) = \sum_{i=1}^n \phi(l_i; F, x_i),$$

where x_i is the i th row of X and l_i is the i th row of L . Alternatively, it can be recovered as

$$(1.7) \quad -\log p(x | F, L) = \sum_{j=1}^p \phi(f_j; L, x_j),$$

in which x_j is the j th column of X , and f_j is the j th row of F . Therefore, when F is fixed, each row of L can be separately optimized by solving a problem of the form (1.4), and when L is fixed, each row of F can be separately optimized by solving a problem of the form (1.4).

Initially this may seem like a sensible strategy, but directly optimizing (1.4) turns out to be difficult to do for numerical reasons: in practice, the entries of B can be very large or very small, resulting in solutions y in which all the entries are either very large or very small. This makes it difficult to devise an algorithm that will work well for all possible input matrices B .

I propose to solve for y indirectly by instead solving

$$(1.8) \quad \begin{aligned} &\text{minimize} && f(t; P, u) \\ &\text{subject to} && t_k \geq 0 \text{ for all } k = 1, \dots, K, \end{aligned}$$

in which the new objective function is

$$(1.9) \quad f(t; P, u) = -\sum_{i=1}^n u_i \log \left(\sum_{k=1}^K p_{ik} t_k \right) + \sum_{k=1}^K t_k,$$

where I've defined

$$\begin{aligned} u_i &= \frac{w_i}{\sum_{i'=1}^n w_{i'}} \\ p_{ik} &= b_{ik} \times \frac{\sum_{i'=1}^n w_{i'}}{\sum_{i'=1}^n b_{i'k}}. \end{aligned}$$

After finding the solution t^* to (1.8), the solution y^* to (1.4) is recovered as

$$(1.10) \quad y_k^* = t_k^* \times \frac{\sum_{i=1}^n w_i}{\sum_{i=1}^n b_{ik}}.$$

The main advantage of solving (1.8) is that the solution is numerically well behaved; in particular, the entries of the solution t^* sum to 1 [2]. A straightforward way to appreciate this result is to compare the complementary slackness condition $\sum_{k=1}^K g_k x_k = 0$ for the original and modified subproblems, where g_k denotes the partial derivative of the objective (either ϕ or f) with respect to the k th co-ordinate (either y_k or t_k). Further, since problem (1.8) is obtained by a simple linear transformation of the variables, any iterate that leads to an improvement in the solution to the modified subproblem will produce an improvement in the solution to the original subproblem, and vice versa.

REFERENCES

- [1] D. M. BLEI, A. Y. NG, AND M. I. JORDAN, *Latent Dirichlet allocation*, Journal of Machine Learning Research, 3 (2003), pp. 993–1022.
- [2] Y. KIM, P. CARBONETTO, M. STEPHENS, AND M. ANITESCU, *A fast algorithm for maximum likelihood estimation of mixture proportions using sequential quadratic programming*, arXiv, 1806.01412 (2019), <https://arxiv.org/abs/1806.01412>.
- [3] D. D. LEE AND H. S. SEUNG, *Algorithms for non-negative matrix factorization*, in Advances in Neural Information Processing Systems 13, 2001, pp. 556–562.